

# 50 High-Quality MCQs on LangChain LCEL & LangGraph

*Instructions: Each question has exactly one correct option. Questions are unique, code- and behavior-focused. No answers are provided.*

1. In LCEL, what is the primary purpose of the pipe (|) operator?
  - A. To merge multiple prompts into one
  - B. To pass the output of one runnable as input to the next
  - C. To execute runnables concurrently
  - D. To store intermediate state globally
2. Which LCEL feature enables composability without immediate execution?
  - A. Static graph compilation
  - B. Lazy runnable evaluation
  - C. Token streaming
  - D. Prompt caching
3. What happens internally when ` `.invoke()` is called on an LCEL chain?
  - A. The chain is converted to JSON
  - B. All runnables execute in sequence with data propagation
  - C. The LLM is re-initialized
  - D. The prompt variables are frozen
4. Which method is recommended for constructing role-based prompts in ChatPromptTemplate?
  - A. from\_template
  - B. from\_prompt
  - C. from\_messages
  - D. from\_schema
5. Why is ChatPromptTemplate preferred over PromptTemplate for chat models?
  - A. It supports tool calling automatically
  - B. It enforces message role separation
  - C. It prevents hallucinations
  - D. It eliminates prompt variables
6. What is the defining characteristic of RunnableLambda?
  - A. It wraps an LLM with memory
  - B. It adapts a Python callable into a runnable
  - C. It executes shell commands
  - D. It creates async graphs
7. How does RunnableLambda handle inputs and outputs?
  - A. Only strings are allowed
  - B. It enforces JSON schemas strictly
  - C. It passes Python objects transparently
  - D. It converts everything to text
8. In an LCEL chain, where is RunnableLambda most commonly used?
  - A. As a replacement for retrievers
  - B. For lightweight data transformation or routing
  - C. To manage long-term memory
  - D. To call external APIs directly
9. What occurs if a RunnableLambda raises an exception during execution?
  - A. The chain retries automatically
  - B. The exception propagates and halts execution

- C. The error is silently ignored
  - D. The output is set to None
10. Which capability allows RunnableLambda to participate in LCEL pipelines?
- A. Runnable interface compliance
  - B. Graph node inheritance
  - C. Token streaming support
  - D. Built-in retry logic
11. What is the core purpose of RunnableParallel?
- A. To batch prompts for an LLM
  - B. To execute multiple runnables on the same input simultaneously
  - C. To merge multiple graphs
  - D. To create async event loops
12. How are outputs from RunnableParallel returned?
- A. As a single concatenated string
  - B. As a dictionary keyed by runnable name
  - C. As a list in execution order
  - D. As separate streams
13. What input does RunnableParallel receive?
- A. Separate inputs for each runnable
  - B. A shared input distributed to all runnables
  - C. Only dictionary inputs
  - D. Only LLM outputs
14. Which scenario best suits RunnableParallel usage?
- A. Multi-step reasoning chains
  - B. Fan-out feature extraction pipelines
  - C. Stateful agent loops
  - D. Prompt engineering experiments
15. How does RunnableParallel differ from LCEL sequential chaining?
- A. It changes prompt formatting
  - B. It avoids shared inputs
  - C. It enables concurrent execution paths
  - D. It enforces strict typing
16. What does LCEL fundamentally abstract away from the developer?
- A. LLM configuration
  - B. Execution orchestration details
  - C. Prompt design
  - D. Tool definitions
17. Which LCEL property enables easy refactoring of chains?
- A. Global state storage
  - B. Functional composition
  - C. Hard-coded execution order
  - D. Implicit retries
18. How does LCEL handle intermediate outputs between runnables?
- A. Via global variables
  - B. Via implicit return-value passing
  - C. Via message history
  - D. Via disk persistence
19. What distinguishes LCEL from traditional imperative LangChain code?

- A. It removes Python entirely
  - B. It emphasizes declarative data flow
  - C. It enforces async-only execution
  - D. It eliminates debugging
20. Which LCEL feature simplifies reuse of partial pipelines?
- A. Chain inheritance
  - B. Runnable composition
  - C. Prompt cloning
  - D. Memory injection
21. In LangGraph, what replaces linear chain execution?
- A. Nested runnables
  - B. Directed graphs with state
  - C. Prompt stacks
  - D. Token queues
22. What is the primary responsibility of a LangGraph node?
- A. Rendering UI components
  - B. Mutating or reading shared state
  - C. Storing embeddings
  - D. Managing LLM tokens
23. How is state represented in LangGraph?
- A. As global variables
  - B. As an explicit typed state object
  - C. As hidden memory buffers
  - D. As prompt templates
24. How does LangGraph decide the next node to execute?
- A. Random selection
  - B. Conditional edges and state evaluation
  - C. Execution timestamps
  - D. Prompt similarity
25. What advantage do graphs provide over chains for agent workflows?
- A. Reduced token usage
  - B. Explicit control flow and branching
  - C. Faster LLM inference
  - D. Simpler syntax
26. What does `create\_react\_agent` abstract away from developers?
- A. Prompt engineering
  - B. Manual ReAct loop implementation
  - C. Tool definitions
  - D. LLM selection
27. Which reasoning pattern does create\_react\_agent implement internally?
- A. Tree-of-Thought
  - B. ReAct
  - C. Chain-of-Verification
  - D. Self-Consistency
28. What inputs are mandatory when creating a ReAct agent in LangGraph?
- A. Only a prompt template
  - B. LLM and tool definitions
  - C. Memory and embeddings
  - D. Dataset and schema

29. How does create\_react\_agent manage iterations?
- A. Fixed number of steps
  - B. Until a termination condition is met
  - C. Until token limit is exceeded
  - D. Until all tools are called
30. What role does state play in create\_react\_agent execution?
- A. It stores UI configuration
  - B. It tracks thoughts, actions, and observations
  - C. It replaces memory modules
  - D. It formats prompts
31. What problem does create\_supervisor\_agent primarily address?
- A. Prompt optimization
  - B. Coordination of multiple agents
  - C. Vector search scaling
  - D. Token streaming
32. How does a supervisor agent differ from a worker agent?
- A. It executes tools directly
  - B. It controls routing and delegation
  - C. It replaces the LLM
  - D. It stores embeddings
33. What architectural pattern best describes a supervisor agent system?
- A. Observer pattern
  - B. Master-worker pattern
  - C. Singleton pattern
  - D. Factory pattern
34. How do worker agents typically communicate results to the supervisor?
- A. Through shared LangGraph state
  - B. Via REST APIs
  - C. Through logs only
  - D. Via prompt templates
35. What happens when multiple worker agents update state?
- A. Updates are ignored
  - B. State transitions follow graph rules
  - C. The system crashes
  - D. Outputs are concatenated
36. How does LangGraph prevent uncontrolled agent execution loops?
- A. Token limits
  - B. Explicit graph termination conditions
  - C. Prompt constraints
  - D. Timeouts only
37. What is the supervisor agent's role in error handling?
- A. Ignore worker failures
  - B. Decide retries, fallback, or termination
  - C. Automatically fix errors
  - D. Log errors only
38. How do conditional edges enhance supervisor-agent graphs?
- A. By reducing token usage
  - B. By enabling dynamic routing decisions

- C. By enforcing static execution
  - D. By removing state
39. Why is `create_supervisor_agent` useful for complex workflows?
- A. It simplifies UI rendering
  - B. It centralizes decision-making logic
  - C. It eliminates the need for tools
  - D. It guarantees perfect outputs
40. How does supervisor-based design improve scalability?
- A. By sharing prompts
  - B. By delegating tasks across specialized agents
  - C. By reducing graph nodes
  - D. By caching LLM responses
41. How can LCEL be used inside LangGraph nodes?
- A. As a replacement for state
  - B. As node execution logic
  - C. As a memory backend
  - D. As a routing engine
42. What benefit does LCEL provide when defining node logic?
- A. Automatic retries
  - B. Clean, composable transformations
  - C. Graph visualization
  - D. Token compression
43. How is RunnableLambda commonly embedded in a LangGraph node?
- A. As a standalone graph
  - B. As the node's callable execution function
  - C. As a memory store
  - D. As a prompt validator
44. What does LangGraph provide that LCEL alone does not?
- A. Prompt templates
  - B. Explicit stateful control flow
  - C. Runnable composition
  - D. Lambda execution
45. How does state mutation differ between LCEL and LangGraph?
- A. LCEL mutates global state
  - B. LangGraph enforces explicit state transitions
  - C. LCEL stores state implicitly
  - D. LangGraph avoids state entirely
46. What must be carefully managed when enabling parallelism in LangGraph?
- A. Prompt length
  - B. State consistency
  - C. Token streaming
  - D. Tool schemas
47. How does LangGraph support iterative reasoning better than LCEL?
- A. By enforcing retries
  - B. By allowing cycles in the graph
  - C. By reducing latency
  - D. By caching outputs
48. Which debugging advantage does LangGraph offer?

- A. Automatic answer validation
  - B. Visibility into state transitions
  - C. Prompt auto-correction
  - D. Reduced hallucinations
49. How does `create_react_agent` leverage LangGraph's execution model?
- A. As a static prompt runner
  - B. As a state-driven reasoning loop
  - C. As a batch inference tool
  - D. As a memory store
50. When combining LCEL with supervisor agents, what is most critical?
- A. Prompt verbosity
  - B. Clear separation of concerns
  - C. Token minimization
  - D. Tool count